

leigqNEWTON

API Reference

Core solver

`leigqNEWTON`

```
help leigqNEWTON
```

leigqNEWTON Left eigenpairs of a quaternionic matrix via a Newton-type solver (stand-alone).

```
[lambda, V, res] = leigqNEWTON(A)
[lambda, V, res, info] = leigqNEWTON(A, Name, Value, ...)
[lambda, V, res, info, lambdaU, VU, resU] = leigqNEWTON(A, Name, Value, ...)
[...] = leigqNEWTON(A, Num, Name, Value, ...) % convenience positional Num
```

Computes quaternionic LEFT eigenpairs (`lambda, V`) of a square quaternion matrix `A`:

$A^*v = \lambda v$, $v \approx 0$,

where `lambda` is a quaternion acting on the LEFT.

This implementation is self-contained and uses only:

- MATLAB built-in quaternion class (Aerospace Toolbox),
- standard MATLAB linear algebra on REAL/COMPLEX embeddings built locally.

Outputs (types and meaning)

`lambda` : K-by-1 quaternion

The K eigenvalue "hits" accepted by the solver (K is controlled by `Num`).
These may contain duplicates (the solver does not enforce uniqueness).

`V` : n-by-K quaternion

Corresponding right eigenvectors as columns ($A^*V(:,k) \approx \lambda(k)V(:,k)$).

`res` : K-by-1 double

Final residual norms reported for each accepted eigenpair.

If `ResidualNormalized=true` (default), `res` is scale-invariant:

$\text{res} = \|A^*v - \lambda v\| / \text{den}(A, \lambda, v)$,

where `den(...)` is a mild scale factor (see `local_residual_den`).

If `ResidualNormalized=false`, `res` is the raw Euclidean norm.

`info` : cell array (only if requested)

`info{1}` is a summary struct; `info{2:end}` are per-trial structs when `InfoLevel='full'`.
In addition, `info{1}.distinct` describes how `lambdaU` was formed (if requested).

`lambdaU` : Ku-by-1 quaternion (only if requested)

A DISTINCT representative set extracted from `lambda` using a tolerance.

Each group of near-identical hits contributes one representative, chosen by:

(i) significantly better residual (if present), otherwise

(ii) "prettiness" (prefers exact integers/zeros over small numerical tails).

By design, `lambdaU` is a subset of the returned `lambda` (representatives are picked from existing hits; no additional cleaning is applied).

`VU` : n-by-Ku quaternion (only if requested)

`resU` : Ku-by-1 double (only if requested)

The vectors/residuals corresponding to `lambdaU`.

Profiles (recommended)

leigqNEWTON is restart-based; success depends strongly on the TRIAL budget and
Newton iterations per trial. SolveProfile presets apply only to options you did

NOT specify explicitly:

```
'fast'      : small budgets; quick scans (may return fewer eigenpairs).
'default'   : balanced; typically succeeds for generic dense matrices.
'reliable'  : larger budgets + auto-extension; aims to return (almost) all eigenpairs.
```

Options (Name,Value) (aliases accepted; unknown options are ignored)

Desired count:

```
'Num'/'K'/'NEigs'/'NEV'      : number of eigenpairs requested (default n)
(positional) Num              : leigqNEWTON(A, 9, ...) sets Num=9
```

Trial budgets / restarts:

```
'Trials'/'Restarts'/'NTrials' : initial trial budget (profile-dependent default)
'TrialsFactor'                : convenience Trials := TrialsFactor*n (if Trials not set)
'MaxTrials'/'MaxRestarts'    : hard cap on total trials (also disables AutoExtend)
'AutoExtend'                  : true/false (profile-dependent default)
'MaxTrialsCap'/'TrialsCap'   : cap used by AutoExtend
'ExtendBy'/'ExtendFactor'    : how much to increase budget when auto-extending
```

Newton / acceptance:

```
'MaxIter'/'MaxIt'           : Newton iterations per trial (profile-dependent default)
'Tol'/'ResTol'                : convergence tolerance (raw residual is used internally)
'Damping'/'Alpha'/'StepSize'  : initial step size alpha in (0,1] (default 1)
'Backtrack'/'LineSearch'     : backtracking line search (default true)
'MinAlpha'/'AlphaMin'         : minimum alpha in backtracking (default 1/64)
```

Distinct representative extraction (for lambdaU):

```
'DistinctTol'/'DistinctTolAbs'/'UniqueTol' : absolute grouping tolerance (default max(1e-6, sqrt(Tol)))
'DistinctTolRel'               : relative grouping tolerance (default 0)
'DistinctResFactor'            : residual "order-of-magnitude" factor (default 10)
```

Note: The distinct-grouping tolerance is intentionally more tolerant than Tol by default.

You still see all raw hits in lambda; lambdaU only groups near-duplicates into representatives.

Reproducibility / logging:

```
'Seed'/'RNG'                 : RNG seed (default [])
'Verbose'/'Disp'               : 0/1 (default 0)
'InfoLevel'/'InfoMode'         : 'summary'|'full' OR numeric 0/1/2 (2='full')
```

Robustness / post-processing:

```
'VerifyZeroNull'/'ZeroNullVerify' : verify Aq*x approx 0 in the zero-eigenvalue pre-pass (default true)
'ZeroNullTol'                   : tolerance for the zero-eig pre-pass (raw units). [] -> auto
'UseNullFallbackLA'/'FallbackLA' : if verification fails, recompute nullspace using real 4n embedding (default false)
'ResidualNormalized'            : if true, res output is scale-invariant (default true)
'RefinerV'                      : recompute eigenvectors for final lambdas via real least-squares/SVD (default false)
```

Initialization / triangular handling:

```
'Lambda0'/'Lam0'/'InitLambda'  : initial lambda guess (scalar quaternion)
'V0'/'X0'/'InitVec'            : initial eigenvector guess (n-by-1 quaternion)
'TriangularInit'/'TriInit'     : if true, deterministic diagonal/basis seeding for (detected) triangular A (default false)
'TriTol'/'TriangularTol'        : triangular detection tolerance (default 0)
'IstTriangular'/'TriangularFlag'/'ForceTriangular' : user override for triangular flag (true/false)
'TriangularShortcut'/'TriShortcut' : 'diag' (default) shortcut for diagonal A; 'off' forces Newton
```

Notes on triangular matrices:

If TriangularInit is true and A is treated as triangular, the solver uses deterministic diagonal/basis seeding for the first n trials. This often leads to very small iteration counts (sometimes 0), which is expected.

To benchmark "generic" random initialization on triangular A, set:

'TriangularInit', false.

Useful one-liners (copy/paste)

% Default (balanced):

```

[lam,V,res] = leigqNEWTON(A);

% Also return distinct representatives (lambdaU):
[lam,V,res,info,lambdaU,VU,resU] = leigqNEWTON(A,'SolveProfile','reliable');
% Add a larger distinct-grouping tolerance explicitly:
[lam,V,res,info,lambdaU,VU,resU] = leigqNEWTON(A,'SolveProfile','reliable','DistinctTolAbs', 1e-5); % or 3e-5, 1e-
% If you want it scale-aware (relative as well)
[lam,V,res,info,lambdaU,VU,resU] = leigqNEWTON(A, 'SolveProfile','reliable','DistinctTolAbs', 1e-5,'DistinctTolRel

% Rule of thumb: Start with DistinctTolAbs = 10*sqrt(Tol) (or 30*sqrt(Tol) for very defective/triangular-ish cas
This affects only deduplication into lambdaU. You still see all raw hits in lam, so being more tolerant here is u

% Request exactly 9 eigenpairs (two equivalent forms):
[lam,V,res] = leigqNEWTON(A, 9, 'SolveProfile','reliable');
[lam,V,res] = leigqNEWTON(A,'Num',9,'SolveProfile','reliable');

% Reproducible run:
[lam,V,res] = leigqNEWTON(A,'Seed',1,'SolveProfile','default');

-----

```

Legacy name mapping

leigqVANILA -> leigqNEWTON

See also: quaternion, parts, null, rank, leigqNEWTON_refine_polish, leigqNEWTON_cert_resMin

Certification

leigqNEWTON_cert_resMin
leigqNEWTON_cert_resPair

help leigqNEWTON_cert_resMin

leigqNEWTON_cert_resMin Minimum-residual certificate for a quaternion LEFT eigenvalue candidate.

```

resMin = leigqNEWTON_cert_resMin(A, lambda)
[resMin, xMin] = leigqNEWTON_cert_resMin(A, lambda)
[resMin, xMin, rMin, info] = leigqNEWTON_cert_resMin(A, lambda, Name,Value,...)

```

Defines the certificate-like quantity

$\text{resMin}(A, \lambda) := \min_{\{||x||_2=1\}} ||A^*x - \lambda x||_2$,
computed via a real/complex embedding and a smallest-singular-value routine.

Inputs

A : n-by-n quaternion (or numeric; interpreted as real quaternion)
lambda : scalar quaternion or K-by-1 quaternion array

Name,Value options (selected)

```

'Method'           : 'svd'|'eigs' (default 'svd') % internal method for sigma_min
'MaxIter'          : iterations for iterative methods (if used)
'Tol'              : tolerance for iterative methods (if used)
'ResidualNormalized' : if true, return scale-invariant resMin (default true)
'UseLambda'        : 'lambda'|'lambdac' (default 'lambda') % choose which lambda vector to use
'LambdaC'          : cleaned lambdas (required if UseLambda='lambdac')

```

Outputs

```

resMin : K-by-1 double (certificate value for each lambda)
xMin   : n-by-K quaternion (minimizers; columns; returned if requested)
rMin   : cell or struct (auxiliary residual vectors/metrics; for diagnostics)
info   : struct       (method statistics; timings; etc.)

```

Useful one-liners

```
r = leigqNewton_cert_resMin(A, lam(1));
[r,x] = leigqNewton_cert_resMin(A, lam, 'ResidualNormalized',true);
```

See also: leigqNewton_cert_resPair, leigqNewton_refine_lambda, leigqNewton

```
help leigqNewton_cert_resPair
```

leigqNewton_cert_resPair Residual certificate for a quaternion LEFT eigenpair (lambda,v).

```
resPair = leigqNewton_cert_resPair(A, lambda, v)
[resPair, rVec] = leigqNewton_cert_resPair(A, lambda, v)
[resPair, rVec, info] = leigqNewton_cert_resPair(A, lambda, v, Name,Value,...)
```

Computes $\|A^*v - \lambda v\|_2$ for one or many eigenpairs. If v is provided as n-by-K (columns), lambda may be scalar or K-by-1.

Inputs

```
A      : n-by-n quaternion (or numeric; interpreted as real quaternion)
lambda : scalar quaternion or K-by-1 quaternion array
v      : n-by-1 or n-by-K quaternion array (eigenvectors in columns)
```

Name,Value options (selected)

```
'NormalizeV'       : true/false (default true). If true, normalize each column of v.
'ResidualNormalized' : true/false (default false). If true, return a scale-invariant residual.
```

Outputs

```
resPair : K-by-1 double      (residual norm(s))
rVec    : cell array         (optional residual vectors; for diagnostics)
info    : struct              (diagnostics; normalization; denominators if used)
```

Useful one-liners

```
r = leigqNewton_cert_resPair(A, lam(1), V(:,1));
r = leigqNewton_cert_resPair(A, lam, V, 'NormalizeV',true);
```

See also: leigqNewton_cert_resMin, leigqNewton_refine_polish, leigqNewton

Refinement

```
leigqNewton_init_vec
leigqNewton_refine_polish
leigqNewton_refine_lambda
leigqNewton_refine_auto
leigqNewton_refine_batch
```

```
help leigqNewton_refine_batch
```

leigqNewton_refine_batch Refine + certify a list of eigenvalue candidates.

```
[lambdaRef, vRef] = leigqNewton_refine_batch(A, lambda0)
[lambdaRef, vRef, cert] = leigqNewton_refine_batch(A, lambda0, Name,Value,...)
```

Batch wrapper that turns coarse candidates lambda0 into refined/certified eigenvalues (and optionally polished eigenvectors). This is the recommended “production” post-processing step after leigqNewton.

Inputs

```
A      : n-by-n quaternion
lambda0 : K-by-1 quaternion candidates (raw or cleaned)
```

Name,Value options (selected)

```
'Mode'       : 'auto' | 'rand+fminsearch' | 'fminsearch' | 'none' (default 'auto')
```

```

'TargetResMin' : target certificate value (default 1e-13)
'DoPolish'     : true/false (default true)

Outputs
lambdaRef : K-by-1 quaternion (refined)
vRef      : n-by-K quaternion (associated vectors; columns)
cert      : struct with fields such as:
             resMin, resPair, lambdaClean, acceptedMask, timings, ...

Useful one-liners
[lamR,vR,cert] = leigqNEWTON_refine_batch(A, lam);
[lamR,vR,cert] = leigqNEWTON_refine_batch(A, lam, 'DoPolish',true,'TargetResMin',1e-14);

See also: leigqNEWTON_refine_auto, leigqNEWTON_refine_lambda, leigqNEWTON_cert_resPair

```

Sphere module

```

leigqNEWTON_sphere_sample
leigqNEWTON_sphere_refine

```

```
help leigqNEWTON_sphere_refine
```

```
leigqNEWTON_sphere_refine Refine/validate sphere candidates and decide sphere-vs-artifact (advanced engine).
```

This routine is the engine behind LEIGQNEWTON_SPHERE_VALIDATE. It is intended for sphere-hunting experiments: it refines sampled eigenvalue candidates, computes residual certificates, and (optionally) verifies a detected sphere by grid tests and refitting. Use LEIGQNEWTON_SPHERE_VALIDATE for a user-facing entry point.

Canonical form (case list + index)

```
[A, lamAll, resAll, lamSamples, resSamples, sph, conf, out] = ...
    leigqNEWTON_sphere_refine(cases, k, Name,Value,...)
```

Convenience overloads

```
leigqNEWTON_sphere_refine(caseStruct, Name,Value,...)
leigqNEWTON_sphere_refine(A, lamAll, lamSamples, Name,Value,...)
```

Inputs

cases / caseStruct Case container (struct, struct array, or cell array). Each case should provide A and lambda candidates.
k Case index (when cases is a list). Use [] for a single case struct.

Positional form inputs

```
A          n-by-n quaternion matrix.
lamAll     M-by-1 quaternion, candidate list.
lamSamples K-by-1 quaternion, DISTINCT samples (may be []).
```

Name,Value options (selected)

```
'RefineArgs'      Cell array of Name,Value passed to LEIGQNEWTON_REFINE_AUTO.
'TargetRes'       Target residual for polishing (default 1e-14).
'VerifySphere'   Enable sphere validation tests (default true).
'SphereTolRel'   Relative tolerance for sphere validation (default 5e-3).
'GridTheta','GridPhi' Sphere grid resolution (defaults 8 and 16).
'RefineGrid'     Refine grid points (default true).
'RefitSphere'   Refit sphere after refinement (default true).
'Verbose'        0/1/2 (default 1).
```

Outputs

```
A, lamAll, lamSamples Refined lists.
resAll, resSamples Certificate-like residual values (via LEIGQNEWTON_CERT_RESMIN).
sph                  Sphere model struct (empty if none/invalid).
```

conf Confidence indicator/struct.
out Detailed diagnostics.

Examples

```
% Validate a detected sphere (positional form):  
[A, lamAll2, resAll, lamS2, resS, sph2] = leigqNewton_sphere_refine(A, lamAll, lamSamples);
```

See also `leigqNewton_sphere_validate`, `leigqNewton_sphere_sample`,
`leigqNewton_refine_auto`, `leigqNewton_cert_resMin`